

Maximum likelihood estimation for regression models

Ciaran Evans

Maximum likelihood estimation and linear regression

Let $(\mathbf{x}_1, Y_1), \dots, (\mathbf{x}_n, Y_n)$ be iid samples from the model

$$Y_i | \mathbf{x}_i \sim N(\mu_i, \sigma^2)$$

$$\mu_i = \mathbf{x}_i^T \boldsymbol{\beta}$$

where the distribution of \mathbf{x}_i does not depend on $\boldsymbol{\beta}$ or σ^2 .

$$L(\boldsymbol{\beta}, \sigma^2 | \mathbf{y}, \mathbf{X}) =$$

Maximum likelihood estimation and logistic regression

Let $(\mathbf{x}_1, Y_1), \dots, (\mathbf{x}_n, Y_n)$ be iid samples from the model

$$Y_i | \mathbf{x}_i \sim \text{Bernoulli}(p_i)$$

$$\log \left(\frac{p_i}{1 - p_i} \right) = \mathbf{x}_i^T \boldsymbol{\beta}$$

where the distribution of \mathbf{x}_i does not depend on $\boldsymbol{\beta}$.

$$L(\boldsymbol{\beta}, \sigma^2 | \mathbf{y}, \mathbf{X}) \propto \prod_{i=1}^n f(Y_i | \mathbf{x}_i, \boldsymbol{\beta}) =$$

Maximizing

$$\ell(\beta | \mathbf{y}, \mathbf{X}) \stackrel{\text{(up to a constant)}}{=} \sum_{i=1}^n \left\{ Y_i \mathbf{x}_i^T \beta - \log(1 + e^{\mathbf{x}_i^T \beta}) \right\}$$

$$\frac{\partial \ell}{\partial \beta} =$$

Score

Definition (score): Let $\mathbf{y} = (Y_1, \dots, Y_n)$ be a sample of n observations from some distribution with parameter vector θ . Let $L(\theta|\mathbf{y})$ be the likelihood function, and $\ell(\theta|\mathbf{y}) = \log L(\theta|\mathbf{y})$ the log-likelihood.

The **score**, which we will denote $U(\theta)$, is the gradient of the log-likelihood with respect to θ :

$$U(\theta) = \frac{\partial}{\partial \theta} \ell(\theta|\mathbf{y}).$$

Example: For logistic regression: $U(\beta) = \mathbf{X}^T(\mathbf{y} - \mathbf{p})$

Question: How would I solve $\mathbf{X}^T(\mathbf{y} - \mathbf{p}) = \mathbf{0}$?

Newton's method

We want to find β^* such that $U(\beta^*) = \mathbf{0}$. Issue: no closed form solution!

Idea: Approximate $U(\beta^*)$ with a first-order Taylor expansion:

$$U(\beta^*) \approx$$

Newton's method

- ▶ Want β^* such that $U(\beta^*) = \mathbf{0}$
- ▶ Begin with initial estimate $\beta^{(k)}$
- ▶ Iterative updates:

$$\beta^{(r+1)} = \beta^{(r)} - (\mathbf{H}(\beta^{(r)}))^{-1} U(\beta^{(r)})$$

The Hessian

$$U(\beta) = \frac{\partial}{\partial \beta} \ell(\beta | \mathbf{y}, \mathbf{X}) = \mathbf{X}^T (\mathbf{y} - \mathbf{p})$$

$$\mathbf{H}(\beta) = \frac{\partial}{\partial \beta} U(\beta) = \frac{\partial}{\partial \beta} \mathbf{X}^T (\mathbf{y} - \mathbf{p})$$

Putting everything together

Want to maximize the log likelihood $\ell(\beta|\mathbf{y}, \mathbf{X})$.